

```

# Install gensim if not already installed
!pip install gensim

"""

STEP 2: Import Libraries

We import:
- gensim to load the Word2Vec model
- numpy for vector math
- pandas for result tables
- matplotlib and PCA for visualization
"""

# to load the pre-trained GoogleNews Word2Vec model
import gensim
from gensim.models import KeyedVectors

# numerical computations
import numpy as np

# table formatting for output
import pandas as pd

# visualization
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA

# computing cosine similarity manually
from sklearn.metrics.pairwise import cosine_similarity

```

```

Collecting gensim
  Downloading gensim-4.4.0-cp312-cp312-manylinux_2_24_x86_64.manylinux_2_28_x86_64.whl.metadata (8.4 kB)
Requirement already satisfied: numpy>=1.18.5 in /usr/local/lib/python3.12/dist-packages (from gensim) (2.0.2)
Requirement already satisfied: scipy>=1.7.0 in /usr/local/lib/python3.12/dist-packages (from gensim) (1.16.3)
Requirement already satisfied: smart_open>=1.8.1 in /usr/local/lib/python3.12/dist-packages (from gensim) (7.5.0)
Requirement already satisfied: wrapt in /usr/local/lib/python3.12/dist-packages (from smart_open>=1.8.1->gensim) (2.1.1)
  Downloading gensim-4.4.0-cp312-cp312-manylinux_2_24_x86_64.manylinux_2_28_x86_64.whl (27.9 MB)
   _____ 27.9/27.9 MB 63.0 MB/s eta 0:00:00
Installing collected packages: gensim
Successfully installed gensim-4.4.0

```

```

# Download the pre-trained GoogleNews-vectors-negative300 model if it doesn't exist
import os
import gensim.downloader as api

MODEL_NAME = "word2vec-google-news-300"

print(f"Downloading and loading {MODEL_NAME} model...")
# Download the GoogleNews-vectors-negative300 model using gensim's API
model = api.load(MODEL_NAME)

print("Model loaded successfully!")

# print model stats
print("Vocabulary Size:", len(model.key_to_index))
print("Vector Size (dimensions):", model.vector_size)

# display a vector for a sample word
sample_word = "king"
print(f"\nVector for '{sample_word}':\n", model[sample_word])

```

```

Downloading and loading word2vec-google-news-300 model...
[=====] 100.0% 1662.8/1662.8MB downloaded
Model loaded successfully!
Vocabulary Size: 3000000
Vector Size (dimensions): 300

Vector for 'king':
 [ 1.25976562e-01  2.97851562e-02  8.60595703e-03  1.39648438e-01
 -2.56347656e-02 -3.61328125e-02  1.11816406e-01 -1.98242188e-01
 5.12695312e-02  3.63281250e-01 -2.42187500e-01 -3.02734375e-01
 -1.77734375e-01 -2.49023438e-02 -1.67968750e-01 -1.69921875e-01
 3.46679688e-02  5.21850586e-03  4.63867188e-02  1.28906250e-01
 1.36718750e-01  1.12792969e-01  5.95703125e-02  1.36718750e-01

```

1.01074219e-01	-1.76757812e-01	-2.51953125e-01	5.98144531e-02
3.41796875e-01	-3.11279297e-02	1.04492188e-01	6.17675781e-02
1.24511719e-01	4.00390625e-01	-3.22265625e-01	8.39843750e-02
3.90625000e-02	5.85937500e-03	7.03125000e-02	1.72851562e-01
1.38671875e-01	-2.31445312e-01	2.83203125e-01	1.42578125e-01
3.41796875e-01	-2.39257812e-02	-1.09863281e-01	3.32031250e-02
-5.46875000e-02	1.53198242e-02	-1.62109375e-01	1.58203125e-01
-2.59765625e-01	2.01416016e-02	-1.63085938e-01	1.35803223e-03
-1.44531250e-01	-5.68847656e-02	4.29687500e-02	-2.46582031e-02
1.85546875e-01	4.47265625e-01	9.58251953e-03	1.31835938e-01
9.86328125e-02	-1.85546875e-01	-1.00097656e-01	-1.33789062e-01
-1.25000000e-01	2.83203125e-01	1.23046875e-01	5.32226562e-02
-1.77734375e-01	8.59375000e-02	-2.18505859e-02	2.05078125e-02
-1.39648438e-01	2.51464844e-02	1.38671875e-01	-1.05468750e-01
1.38671875e-01	8.88671875e-02	-7.51953125e-02	-2.13623047e-02
1.72851562e-01	4.63867188e-02	-2.65625000e-01	8.91113281e-03
1.49414062e-01	3.78417969e-02	2.38281250e-01	-1.24511719e-01
-2.17773438e-01	-1.81640625e-01	2.97851562e-02	5.71289062e-02
-2.89306641e-02	1.24511719e-02	9.66796875e-02	-2.31445312e-01
5.81054688e-02	6.68945312e-02	7.08007812e-02	-3.08593750e-01
-2.14843750e-01	1.45507812e-01	-4.27734375e-01	-9.39941406e-03
1.54296875e-01	-7.66601562e-02	2.89062500e-01	2.77343750e-01
-4.86373901e-04	-1.36718750e-01	3.24218750e-01	-2.46093750e-01
-3.03649902e-03	-2.11914062e-01	1.25000000e-01	2.69531250e-01
2.04101562e-01	8.25195312e-02	-2.01171875e-01	-1.60156250e-01
-3.78417969e-02	-1.20117188e-01	1.15234375e-01	-4.10156250e-02
-3.95507812e-02	-8.98437500e-02	6.34765625e-03	2.03125000e-01
1.86523438e-01	2.73437500e-01	6.29882812e-02	1.41601562e-01
-9.81445312e-02	1.38671875e-01	1.82617188e-01	1.73828125e-01
1.73828125e-01	-2.37304688e-01	1.78710938e-01	6.34765625e-02
2.36328125e-01	-2.08984375e-01	8.74023438e-02	-1.66015625e-01
-7.91015625e-02	2.43164062e-01	-8.88671875e-02	1.26953125e-01
-2.16796875e-01	-1.73828125e-01	-3.59375000e-01	-8.25195312e-02
-6.49414062e-02	5.07812500e-02	1.35742188e-01	-7.47070312e-02
-1.64062500e-01	1.15356445e-02	4.45312500e-01	-2.15820312e-01
-1.11328125e-01	-1.92382812e-01	1.70898438e-01	-1.25000000e-01
2.65502930e-03	1.92382812e-01	-1.74804688e-01	1.39648438e-01
2.92968750e-01	1.13281250e-01	5.95703125e-02	-6.39648438e-02
9.96093750e-02	-2.72216797e-02	1.96533203e-02	4.27246094e-02
-2.46093750e-01	6.39648438e-02	-2.25585938e-01	-1.68945312e-01
2.89916992e-03	8.20312500e-02	3.41796875e-01	4.32128906e-02
1.32812500e-01	1.42578125e-01	7.61718750e-02	5.98144531e-02
-1.19140625e-01	2.74658203e-03	-6.29882812e-02	-2.72216797e-02
-4.82177734e-03	-8.20312500e-02	-2.49023438e-02	-4.00390625e-01

```
"""
STEP 4: Compute Similarity for Word Pairs
```

We compute cosine similarity: values range from -1 (opposite) to +1 (very similar).

```
"""
```

```
word_pairs = [
    ("doctor", "nurse"),
    ("cat", "dog"),
    ("car", "bus"),
    ("king", "queen"),
    ("apple", "car"),
    ("teacher", "student"),
    ("india", "china"),
    ("hospital", "doctor"),
    ("computer", "laptop"),
    ("river", "ocean")
]

results = []

for w1, w2 in word_pairs:
    score = model.similarity(w1, w2)
    results.append([w1, w2, score])
    print(f"{w1} - {w2}: {score:.4f}")

# convert to table
df = pd.DataFrame(results, columns=["Word1", "Word2", "CosineSimilarity"])
df
```

```

doctor - nurse: 0.6320
cat - dog: 0.7609
car - bus: 0.4693
king - queen: 0.6511
apple - car: 0.1283
teacher - student: 0.6301
india - china: 0.3533
hospital - doctor: 0.5143
computer - laptop: 0.6640
river - ocean: 0.4772

```

	Word1	Word2	CosineSimilarity	grid icon
	Word1	Word2	CosineSimilarity	edit icon
0	doctor	nurse	0.631952	
1	cat	dog	0.760946	
2	car	bus	0.469337	
3	king	queen	0.651096	
4	apple	car	0.128307	
5	teacher	student	0.630137	
6	india	china	0.353308	
7	hospital	doctor	0.514324	
8	computer	laptop	0.664049	
9	river	ocean	0.477181	

Next steps: [Generate code with df](#) [New interactive sheet](#)

"""  
STEP 5: Find Nearest Neighbors

For each chosen word, print the top 5 most similar words.  
"""

```

words_to_check = ["king", "university", "hospital", "computer", "india"]

for w in words_to_check:
    print(f"\nTop 5 similar to '{w}':")
    for neighbor, score in model.most_similar(w, topn=5):
        print(f"{neighbor} ({score:.4f})")

```

Top 5 similar to 'king':  
kings (0.7138)  
queen (0.6511)  
monarch (0.6413)  
crown\_prince (0.6204)  
prince (0.6160)

Top 5 similar to 'university':  
universities (0.7004)  
faculty (0.6781)  
university (0.6758)  
undergraduate (0.6587)  
univeristy (0.6585)

Top 5 similar to 'hospital':  
Hospital (0.7932)  
hopsital (0.7784)  
hosptial (0.7582)  
hospitals (0.7213)  
intensive\_care (0.7206)

Top 5 similar to 'computer':  
computers (0.7979)  
laptop (0.6640)  
laptop\_computer (0.6549)  
Computer (0.6473)  
com\_puter (0.6082)

Top 5 similar to 'india':  
indian (0.6967)  
usa (0.6836)  
pakistan (0.6815)  
chennai (0.6676)

```
america (0.6589)
```

```
"""
STEP 6: Word Analogies
```

```
Word analogies use vector arithmetic: vector relationships encode meaning.
```

```
"""
# Example: king - man + woman ≈ queen
print("king - man + woman =", model.most_similar(positive=["king", "woman"], negative=["man"], topn=1))

# Example: paris - france + india
print("paris - france + india =", model.most_similar(positive=["paris", "india"], negative=["france"], topn=1))

# Example: doctor analogies
print("doctor - man + woman =", model.most_similar(positive=["doctor", "woman"], negative=["man"], topn=1))
```

```
king - man + woman = [('queen', 0.7118193507194519)]
paris - france + india = [('chennai', 0.5442505478858948)]
doctor - man + woman = [('gynecologist', 0.7093892097473145)]
```

```
"""
STEP 7: Manual Cosine Similarity
```

```
Compute similarity between two vectors using sklearn.
```

```
"""
vec1 = model["doctor"].reshape(1, -1)
vec2 = model["nurse"].reshape(1, -1)

manual_score = cosine_similarity(vec1, vec2)[0][0]
print("Manual cosine similarity (doctor, nurse):", manual_score)
```

```
Manual cosine similarity (doctor, nurse): 0.6319524
```

```
"""
STEP 8: Visualize Embeddings with PCA
```

```
We will plot selected words in 2D space using PCA.
```

```
"""
words_for_plot = [
    "king", "queen", "man", "woman",
    "paris", "france", "india", "delhi",
    "doctor", "nurse", "hospital",
    "computer", "laptop", "software",
    "cat", "dog", "lion", "tiger",
    "apple", "banana"
]

vectors = np.array([model[w] for w in words_for_plot])

pca = PCA(n_components=2)
reduced = pca.fit_transform(vectors)

plt.figure(figsize=(12, 8))
for i, word in enumerate(words_for_plot):
    x, y = reduced[i]
    plt.scatter(x, y)
    plt.text(x + 0.02, y + 0.02, word)
plt.title("Word2Vec 300D Embeddings Visualized (PCA)")
plt.xlabel("component 1")
plt.ylabel("component 2")
plt.grid()
plt.show()
```

Word2Vec 300D Embeddings Visualized (PCA)

